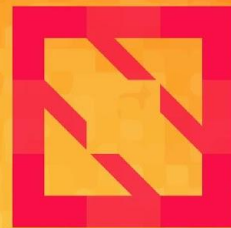




KubeCon



CloudNativeCon

North America 2019





KubeCon



CloudNativeCon

North America 2019

Running High Performance User-space Packet Processing Apps in Kubernetes

Abdul Halim, Intel

Peng Liu, Red Hat



Legal Notices and Disclaimers



KubeCon



CloudNativeCon

North America 2019

- Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Performance varies depending on system configuration **No product or component can be absolutely secure**. Check with your system manufacturer or retailer or learn more at [intel.com](https://www.intel.com).
- Intel, the Intel logo are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries.
- *Other names and brands may be claimed as the property of others.
- © Intel Corporation



Topics



KubeCon



CloudNativeCon

North America 2019

- Motivation
- Challenges and overheads
- Utilizing the NICs line-rates
- DPDK
- Sample deployment
- Capabilities needed in K8s
- SR-IOV networking in K8s
- SR-IOV network operator
- Demo

About us



KubeCon



CloudNativeCon

North America 2019

Abdul Halim

Intel

- Cloud Software Engineer at **Network Control and Logic Group**, Intel
- Enabling high-performance networking solution for NFV with Kubernetes



@ahalim-intel

Peng Liu

Red Hat

- NFV Partner Engineer at **CTO Office** of RedHat
- Working on enhance open source softwares for NFV use cases.



@pliurh

Motivation



KubeCon



CloudNativeCon

North America 2019

The 5G network is expected to enable fully mobile and connected society with the vision of providing:

- Greater throughput
- Lower latency
- Ultra-high reliability
- Much higher connectivity density, and
- Higher user mobility



➤ Demands increasing network speed

5G depends on Kubernetes in the cloud!

Challenges



KubeCon



CloudNativeCon

North America 2019

Adapter	Capability
Intel® Ethernet Network Adapter XXV710-DA2	25/10/1GbE
Intel® Ethernet Converged Network Adapter XL710-QDA1	40/10GbE
Intel® Ethernet Controller E810-CAM2/CAM1	100/50/25/10/1GbE

NIC Capability (Line-rate)	Payloads (Bytes)	Packets per seconds	Packets arrival rate(ns)
10 Gbits/s	1500	812.74Kpps	1230.4
	46	14.88Mpps	67.2
40 Gbits/s	1500	3.25Mpps	307.6
	46	59.52Mpps	16.8
100 Gbits/s	1500	8.12Mpps	123.04
	46	148.80Mpps	6.72

- The higher the packet rate, the lower the time to process it
- To achieve “*zero-packet-loss*”, time spent in network stack must be \leq packets arrival rate

Overheads in Kernel



KubeCon



CloudNativeCon

North America 2019

- For the smallest frame (84 bytes) there is **67.2 ns** to process a packet given 10Gbit/s link
- Which is:
 - 201 CPU cycles @ 3GHz

Ref: [Brouer, J.D, 2015 - Network stack challenges at increasing speeds](#)

Overheads in Kernel



KubeCon



CloudNativeCon

North America 2019

➤ A lot of these cycles could be spent on:

- Context switching
- System calls
- Interrupt handling
- Lock and unlock
- Data copy
- Cache misses

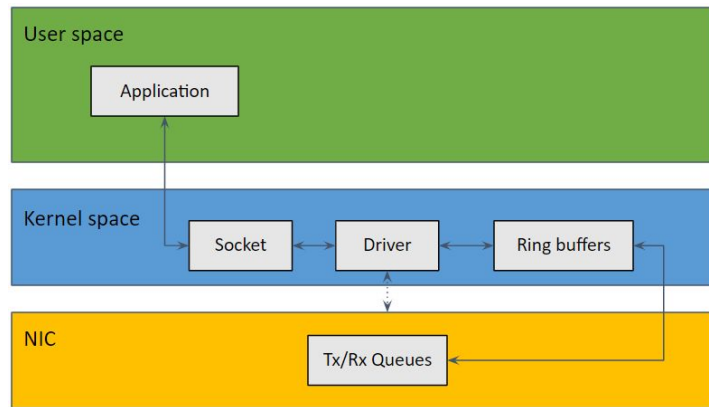


Fig: Packet processing in Linux kernel

Utilizing full line-rates



KubeCon



CloudNativeCon

North America 2019

- By adding performance optimization in Kernel stack
- Network stack bypass solution
 - DPDK - Packet processing in user space
 - RDMA
 - Programmable packet processing with XDP



- ***DPDK (Data Plane Development Kit)*** is a framework (under the Linux Foundation) comprised of various userspace libraries and drivers for fast packet processing

How DPDK does it



KubeCon



CloudNativeCon

North America 2019

➤ DPDK runs in user space

- Dedicated processors -> no context switching
- Dedicated network I/O
- Hugepages -> no swap, TLB
- UIO -> No copy from Kernel
- Polling -> No interrupt overhead
- Lockless synchronization
- Batch packets handling

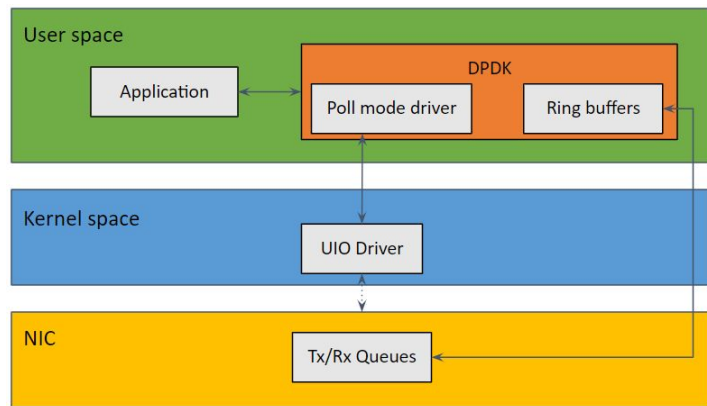


Fig: Packet processing in user space with DPDK

Ref: [Brouer, J.D, 2015 - Network stack challenges at increasing speeds](#)

- But DPDK comes with caveats
- Guaranteed performance requires optimal resource allocation & tuning
 - Exclusive CPU cores
 - Direct device assignments
 - Huge Page memory
 - NUMA alignments
- These are the key requirements in Kubernetes for DPDK apps
 - Challenge here is ***resource allocation & management***

This session about



KubeCon



CloudNativeCon

North America 2019

Capabilities needed in Kubernetes to orchestrate resource
critical DPDK apps and how to do it

Sample deployment



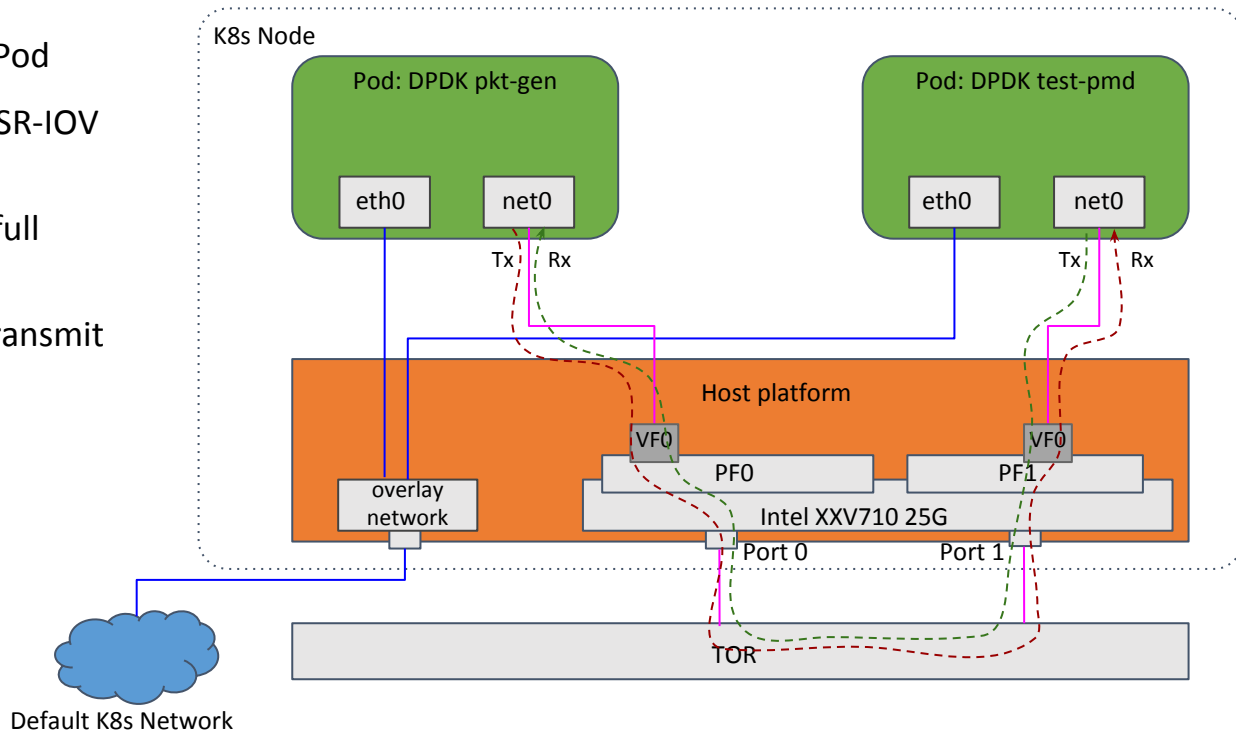
KubeCon



CloudNativeCon

North America 2019

- 2 DPDK apps running in K8s Pod
- Send/Receive packets using SR-IOV VFs
- Pkt-gen transmit packets at full line-rate
- Test-pmd acts as l2fwd, re-transmit packets to the src addr



Platform overview



KubeCon



CloudNativeCon

North America 2019



Intel® Server Board S2600WFR (Wolf Pass Refresh)



Intel® Ethernet Network Adapter
XXV710-DA2 2x25G

Feature	S2600WF0R
Platform	Intel® Server Board S2600WFR (Wolf Pass Refresh)
Processor	2 x Intel® Xeon® Platinum Cascade Lake SP 8280M Processor 28 Cores @ 2.70GHz 38.5MB L3 Cache
Chipset	Intel® C624
Memory	24 DDR4 RDIMM/LRDIMMs, 2 SPC, 12x channels/system
	2666 MT/s @ 2DPC, 72 GB total
PCIe*	Up to 8 PCIe* slots via 3 Risers, One x8 PCIe Gen 3 SAS Mezz Module
NIC (PCIe addons)	Intel® Ethernet Network Adapter XXV710-DA2 2x25G
BIOS	Vendor: Intel Corporation
	Version: SE5C620.86B.0D.01.0321.011120191026
	Release Date: 01/11/2019

Running a DPDK app



KubeCon



CloudNativeCon

North America 2019

Exclusive CPUs

NUMA aligned memory allocation

```
# testpmd -l 3,31 -n 4 -w 0000:18:02.1 -- -i --portmask=0x1 --nb-cores=2 --numa
```

Device assignment

```
EAL: Detected 56 lcore(s)
EAL: Detected 2 NUMA nodes
EAL: Multi-process socket /var/run/dpdk/rte/mp_socket
EAL: Selected IOVA mode 'VA'
EAL: No free hugepages reported in hugepages-2048kB
EAL: Probing VFIO support...
EAL: VFIO support initialized
EAL: PCI device 0000:18:02.1 on NUMA socket 1
EAL: probe driver: 8086:154c net_i40e_vf
EAL: using IOMMU type 1 (Type 1)
Auto-start selected
Set macswap packet forwarding mode
testpmd: create a new mbuf pool <mbuf_pool_socket_0>: n=155456, size=2176, socket=1
Configuring Port 0 (socket 1)
Port 0: BA:AA:7A:88:F2:44
Checking link statuses...
Done
```



Required K8s Capabilities



KubeCon



CloudNativeCon

North America 2019

- Exclusive CPUs
 - Native CPU manager (v1.8)
- Huge Page memory
 - Native Hugepage support (v1.8)
- Direct device assignment
 - Device Plugins (v1.8)
- Resource NUMA alignments
 - Topology manager (v1.16 alpha)

Required other components



KubeCon



CloudNativeCon

North America 2019

- SR-IOV network resource management
 - github.com/intel/sriov-network-device-plugin
- SR-IOV network interface configuration
 - github.com/intel/multus-cni
 - github.com/intel/sriov-cni
- Managing configuration and deployment
 - github.com/openshift/sriov-network-operator
- Pod resource parameters
 - github.com/openshift/app-netutil
- SR-IOV network resource injector
 - github.com/intel/network-resources-injector

SW Configuration



KubeCon



CloudNativeCon

North America 2019

SW	Version
Kubernetes	v1.16
Openshift	v4.3
Host OS: RedHat CoreOS	v4.3
SRIOV CNI Plugin	v2.1
SRIOV Network Device Plugin	v.3.0
SRIOV Network Operator	v4.3
Multus	v3.3
DPDK	v19.08
pktgen-dpdk	v19.08

➤ Getting exclusive CPUs

- Run containers in Pod with “**Guaranteed**” QoS class
- Request exclusive CPU cores
 - Run Kubelet with:
--cpu-manager-policy=static
 - Add equal integer values for CPU in both `requests` & `limits`

Guaranteed Pod:

```
spec:
  containers:
  - name: nginx
    image: nginx
    resources:
      limits:
        memory: "200Mi"
        cpu: "2"
      requests:
        memory: "200Mi"
        cpu: "2"
```

- Enable resource NUMA alignment
 - Available in K8s (v1.16 alpha)
 - Works on Nodes with the static CPU Manager Policy
 - Works on Pods in the **Guaranteed** QoS class
 - Enabling Topology manager
 - Run Kubelet with:
--feature-gates="TopologyManager=true"
 - And ***--topology-manager-policy=restricted***
OR ***single-numa-node***

➤ Requesting Huge Page memory

- Nodes must pre-allocate huge pages
- A node may only pre-allocate a single size
- Add huge pages as Volumes in containers

Guaranteed Pod:

```
spec:
  containers:
  - name: nginx
    image: nginx
    resources:
      limits:
        hugepages-2Mi: 2G
      requests:
        hugepages-2Mi: 2G
    volumeMounts:
    - mountPath: /hugepages
  volumes:
  - name: hugepage
    emptyDir:
      medium: HugePages
```

SR-IOV networking in K8s



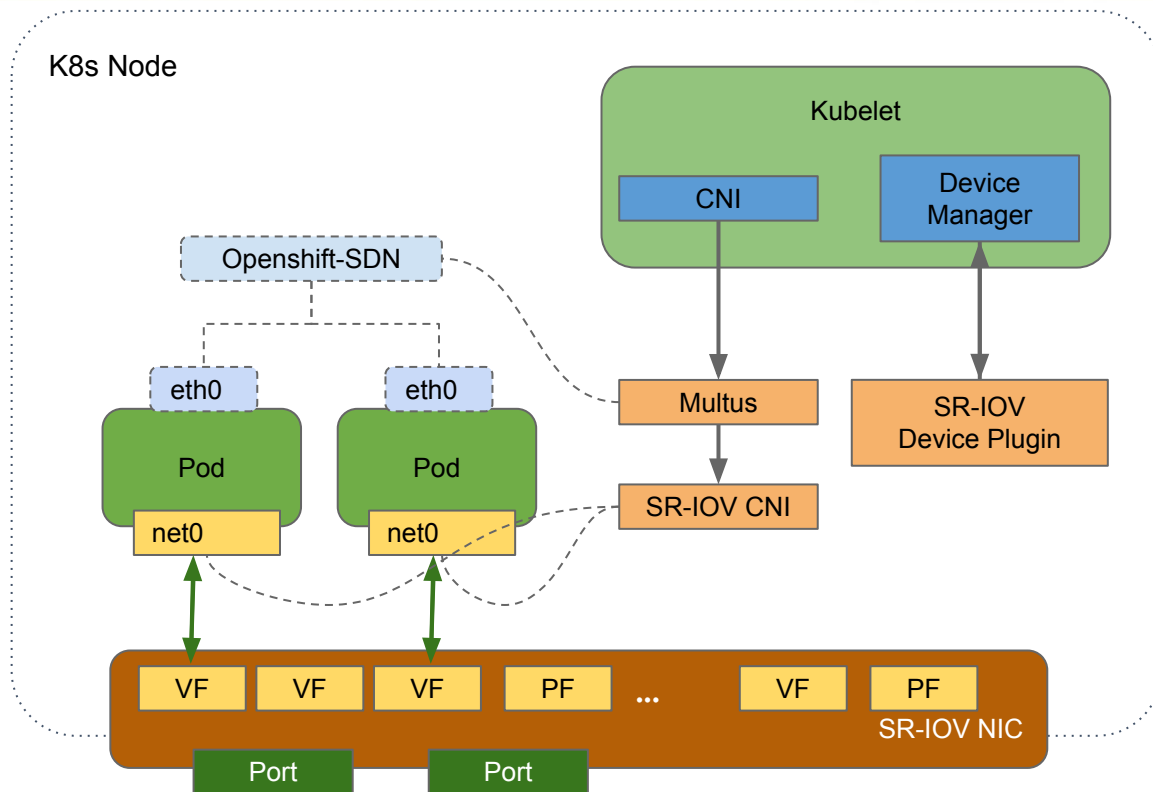
KubeCon



CloudNativeCon

North America 2019

- SR-IOV device plugin
 - Discovery & advertising of SR-IOV network resources
- SR-IOV CNI
 - Configure pod interface
- Multus
 - Adds SR-IOV VF as an additional Pod interface
 - Retrieves Pod device info



SR-IOV networking in K8s



KubeCon



CloudNativeCon

North America 2019

Kube-apiserver

ConfigMap:

```
{
  "resourceList": [
    {
      "resourceName": "nic1",
      "selectors": {
        "vendors": ["8086"],
        "drivers": ["vfio-pci"],
        "pfNames": ["enps803f0"]
      }
    },
    {
      "resourceName": "nic2",
      "selectors": {
        "vendors": ["8086"],
        "drivers": ["vfio-pci"],
        "pfNames": ["enps803f1"]
      }
    }
  ]
}
```

K8s Node

Kubelet

Device
Manager

SR-IOV
Device Plugin

<< discover devices >>

VF

VF

VF

PF

...

VF

PF

SR-IOV NIC

Port

Port

SR-IOV networking in K8s



KubeCon



CloudNativeCon

North America 2019

Kube-apiserver

K8s Node

<< update node status >>

Kubelet

Device
Manager

<< register devices >>

SR-IOV
Device Plugin

SR-IOV NIC

VF

VF

VF

PF

...

VF

PF

Port

Port

Node status:

```
Name: k8s-node1.ir.intel.com
Capacity:
  cpu: 8
  ephemeral-storage: 184447308Ki
  hugepages-1Gi: 0
  hugepages-2Mi: 8Gi
  intel.com/nic1: 4
  intel.com/nic2: 4
  memory: 16371628Ki
  pods: 110
Allocatable:
  cpu: 8
  ephemeral-storage: 169986638772
  hugepages-1Gi: 0
  hugepages-2Mi: 8Gi
  intel.com/nic1: 4
  intel.com/nic2: 4
  memory: 7880620Ki
  pods: 110
```



SR-IOV networking in K8s



KubeCon



CloudNativeCon

North America 2019

Node status

```
Name: k8s-node1.ir.intel.com
Capacity:
  cpu: 8
  ephemeral-storage: 184447308Ki
  hugepages-1Gi: 0
  hugepages-2Mi: 8Gi
  intel.com/nic1: 4
  intel.com/nic2: 4
  memory: 16371628Ki
  pods: 1k
Allocatable:
  cpu: 8
  ephemeral-storage: 169986638772
  hugepages-1Gi: 0
  hugepages-2Mi: 8Gi
  intel.com/nic1: 4
  intel.com/nic2: 4
  memory: 7880620Ki
  pods: 110
```

Net-attach CRD

```
apiVersion: "k8s.cni.cncf.io/v1"
kind: NetworkAttachmentDefinition
metadata:
  name: sriov-net
  annotations:
    k8s.v1.cni.cncf.io/resourceName:
      intel.com/nic1
spec:
  config: '{
    "type": "sriov",
    "name": "sriov-network"
  }'
```

Pod Specs

```
apiVersion: v1
kind: Pod
metadata:
  name: testpod
  annotations:
    k8s.v1.cni.cncf.io/networks:
      openshift-sdn, sriov-net
spec:
  containers:
    - name: appcntr1
      resources:
        requests:
          intel.com/nic1: 1
        limits:
          intel.com/nic1: 1
```



SR-IOV Network Operator



KubeCon



CloudNativeCon

North America 2019

Simplifying SR-IOV networking in K8s with **SR-IOV network operator**



SR-IOV Network Operator



KubeCon



CloudNativeCon

North America 2019

- SR-IOV Network Operator is a tool to hide the complexity of the adopting SR-IOV in K8S. It automates
 - Life-cycle management of SR-IOV software components (SR-IOV CNI plugin, SR-IOV network device plugin, network resource injector ...)
 - The configuration management of the SR-IOV software components.
 - The SR-IOV network device configuration
 - Hardware discovery
 - Device configuration
 - Kernel driver management
 - **NetworkAttachmentDefinition** CR generation
- Built for OpenShift, can also work with vanilla K8S



SR-IOV Network Operator



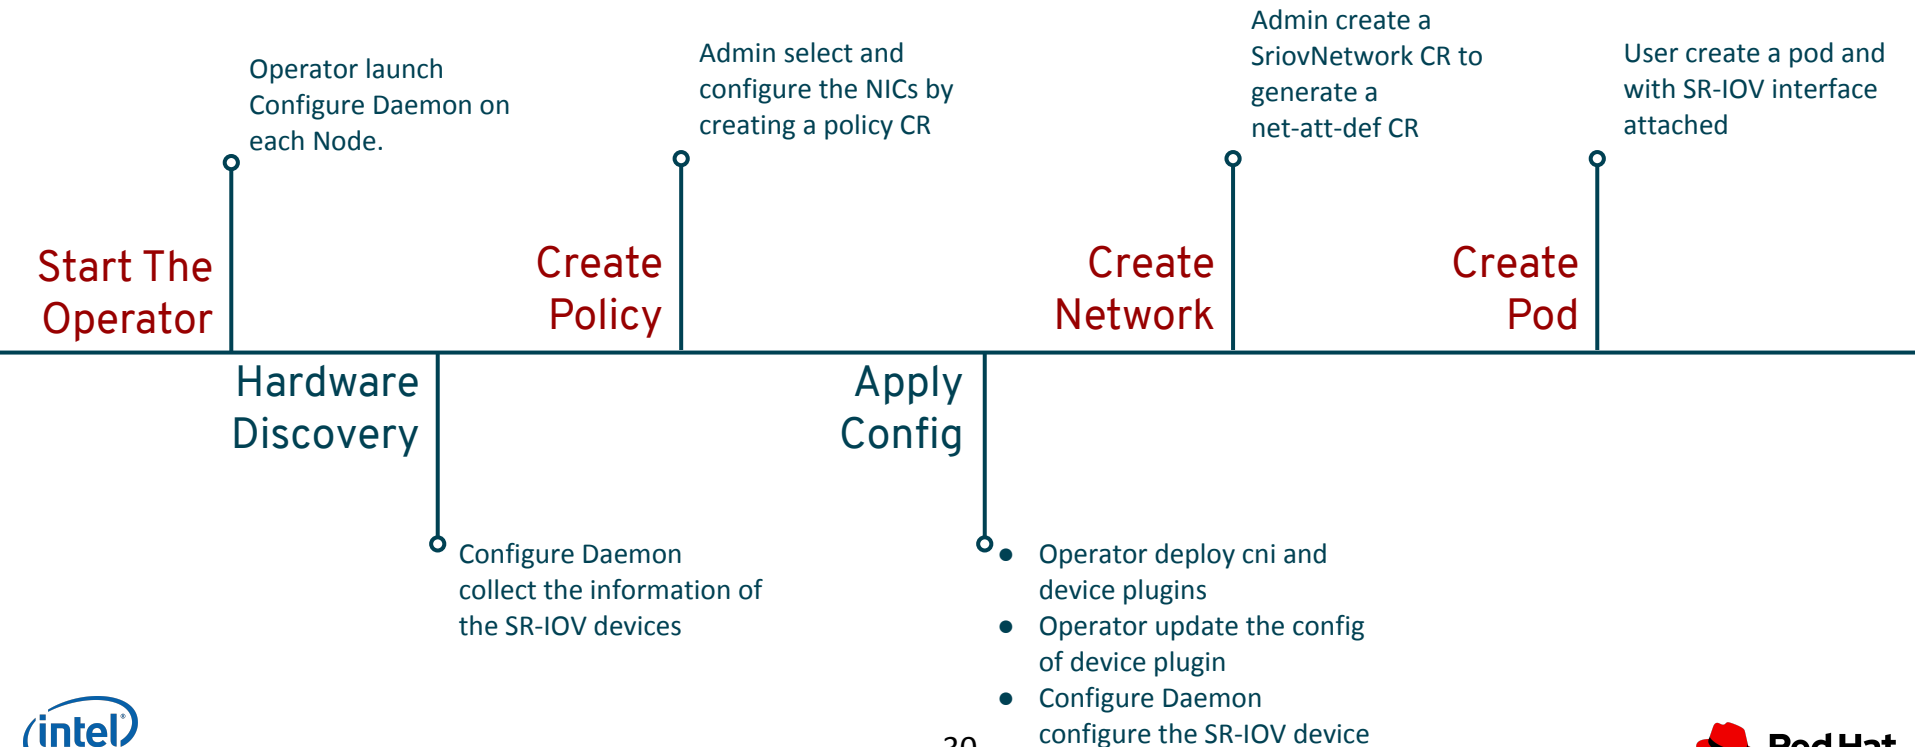
KubeCon



CloudNativeCon

North America 2019

Operator Workflow



SR-IOV Network Operator



KubeCon




CloudNativeCon

North America 2019

SriovNetwork

```
apiVersion:
sriovnetwork.openshift.io/v1
kind: SriovNetwork
metadata:
  name: sriov-net
  namespace: sriov-network-operator
spec:
  networkNamespace: default
  ipam: |
    {
      "type": "host-local",
      "subnet": "10.56.217.0/24",
      "rangeStart": "10.56.217.171",
      "rangeEnd": "10.56.217.181",
    }
  resourceName: nic1
```



Net-attach CRD

```
apiVersion: "k8s.cni.cncf.io/v1"
kind: NetworkAttachmentDefinition
metadata:
  name: sriov-net
  namespace: default
  annotations:
    k8s.v1.cni.cncf.io/resourceName:
intel.com/nic1
spec:
  config: '{
    "type": "sriov",
    "name": "sriov-network",
    "ipam": {
      "type": "host-local",
      "subnet": "10.56.217.0/24",
      "rangeStart": "10.56.217.171",
      "rangeEnd": "10.56.217.181",
    }
  }'
```

SR-IOV Network Operator



KubeCon



CloudNativeCon

North America 2019

SriovNetworkNodePolicy

```
apiVersion: sriovnetwork.openshift.io/v1
kind: SriovNetworkNodePolicy
metadata:
  name: policy-1
  namespace: sriov-network-operator
spec:
  deviceType: vfio-pci
  numVfs: 4
  mtu: 1500
  Priority: 90
  resourceName: nic1
  nicSelector:
    rootDevices:
      - 0000:86:00.1
    vendor: "8086"
    pfName: "enps803f0"
  nodeSelector:
    feature.node.kubernetes.io/network-sriov.capable: "true"
```

Config options for selected VFs:

- Number of VFs
- MTU
- Kernel driver, either 'netdevice' or 'vfio-pci'

Select NICs which need to be configured, and generate SRIOV device plugin config

```
{
  "resourceList": [
    {
      "resourceName": "nic1",
      "selectors": {
        "vendors": ["8086"],
        "drivers": ["vfio-pci"],
        "pfNames": ["enps803f0"]
      }
    }
  ]
}
```

Select Nodes to be configured



SR-IOV Network Operator



KubeCon



CloudNativeCon

North America 2019

SriovNetworkNodeState

```
apiVersion: sriovnetwork.openshift.io/v1
kind: SriovNetworkNodeState
metadata:
```

```
  name: worker-node-1
  namespace: sriov-network-operator
```

```
spec:
```

```
  interfaces:
  - deviceType: vfio-pci
    mtu: 1500
    numVfs: 4
    pciAddress: 0000:86:00.0
```

```
status:
```

```
  interfaces:
  - deviceID: "1583"
    driver: i40e
    mtu: 1500
    numVfs: 4
    pciAddress: 0000:86:00.0
    maxVfs: 64
    vendor: "8086"
    Vfs:
    - deviceID: 154c
      driver: vfio-pci
      pciAddress: 0000:86:02.0
      vendor: "8086"
```

...

Maintained and updated by the operator,
one instance for each worker node

The desired status of the interfaces

The current status of all the interfaces





KubeCon



CloudNativeCon

North America 2019

Demo

- Video: <https://youtu.be/scp2WV5M3TI>
- Sample manifests: <https://github.com/pliurh/Kubecon2019-DEMO>



What needs attention



KubeCon



CloudNativeCon

North America 2019

- Isolated CPU support
 - Kubelet unable to manage isolated CPU cores
 - Multiple OOT solution exist
 - None of these are well integrated with native CPU manager

- Topology aware scheduling
 - K8s default scheduler unaware of resource Topology information
 - May results in Pod scheduled in non-viable nodes

Summary



KubeCon



CloudNativeCon

North America 2019

- A lot of work has been done in the community to support Telco and 5G use-cases in Kubernetes
- Kubernetes is “**5G Ready**”
- Some areas still need attention and wider community collaboration

References



KubeCon



CloudNativeCon

North America 2019

- <http://core.dpdk.org/perf-reports/>
- <https://github.com/intel/multus-cni>
- <https://github.com/intel/network-resources-injector>
- <https://github.com/intel/sriov-cni/>
- <https://github.com/intel/sriov-network-device-plugin>
- <https://github.com/openshift/app-netutil>
- <https://github.com/openshift/sriov-network-operator>
- <https://github.com/pktgen/Pktgen-DPDK>
- <https://kubernetes.io/docs/concepts/configuration/manage-compute-resources-container/>
- <https://kubernetes.io/docs/tasks/administer-cluster/cpu-management-policies/>
- <https://kubernetes.io/docs/tasks/administer-cluster/topology-manager/>
- <https://kubernetes.io/docs/tasks/manage-hugepages/scheduling-hugepages/>
- <https://www.dpdk.org/>
- <https://www.intel.com/content/www/us/en/products/servers/server-chassis-systems/server-board-s2600wf-systems.html>

Additional Resources



KubeCon



CloudNativeCon

North America 2019

Please visit Intel® Network Builders site for Bare-metal Containers
Experience Kits



<https://networkbuilders.intel.com/network-technologies/container-experience-kits>

